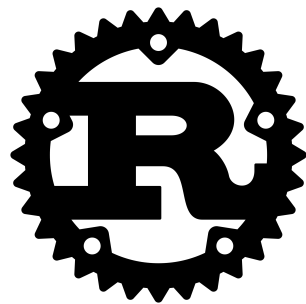


# Introduction to Rust

Ian McCormack





Why would you choose one programming language over another?

**What is systems programming?**

# Rust is popular and widely used in production.

Chosen as **the “most loved” language** in StackOverflow’s annual developer survey for **the last eight years.**




Total

	Energy (J)
(c) C	1.00
(c) Rust	1.03
(c) C++	1.34
(c) Ada	1.70
(v) Java	1.98
(c) Pascal	2.14
(c) Chapel	2.18
(v) Lisp	2.27
(c) Ocaml	2.40
(c) Fortran	2.52
(c) Swift	2.79
(c) Haskell	3.10
(v) C#	3.14
(c) Go	3.23
(i) Dart	3.83
(v) F#	4.13
(i) JavaScript	4.45
(v) Racket	7.91
(i) TypeScript	21.50
(i) Hack	24.02
(i) PHP	29.30
(v) Erlang	42.23
(i) Lua	45.98
(i) Jruby	46.54
(i) Ruby	69.91
(i) Python	75.88
(i) Perl	79.58

	Time (ms)
(c) C	1.00
(c) Rust	1.04
(c) C++	1.56
(c) Ada	1.85
(v) Java	1.89
(c) Chapel	2.14
(c) Go	2.83
(c) Pascal	3.02
(c) Ocaml	3.09
(v) C#	3.14
(v) Lisp	3.40
(c) Haskell	3.55
(c) Swift	4.20
(c) Fortran	4.20
(v) F#	6.30
(i) JavaScript	6.52
(i) Dart	6.67
(v) Racket	11.27
(i) Hack	26.99
(i) PHP	27.64
(v) Erlang	36.71
(i) Jruby	43.44
(i) TypeScript	46.20
(i) Ruby	59.34
(i) Perl	65.79
(i) Python	71.90
(i) Lua	82.91

	Mb
(c) Pascal	1.00
(c) Go	1.05
(c) C	1.17
(c) Fortran	1.24
(c) C++	1.34
(c) Ada	1.47
(c) Rust	1.54
(v) Lisp	1.92
(c) Haskell	2.45
(i) PHP	2.57
(c) Swift	2.71
(i) Python	2.80
(c) Ocaml	2.82
(v) C#	2.85
(i) Hack	3.34
(v) Racket	3.52
(i) Ruby	3.97
(c) Chapel	4.00
(v) F#	4.25
(i) JavaScript	4.59
(i) TypeScript	4.69
(v) Java	6.01
(i) Perl	6.62
(i) Lua	6.72
(v) Erlang	7.20
(i) Haskell	
(i) Java	

Pereira et al., 2021



# What is a memory error?

“...an object accessed using a pointer expression is different from the one intended.”

— van der Veen et al., 2012

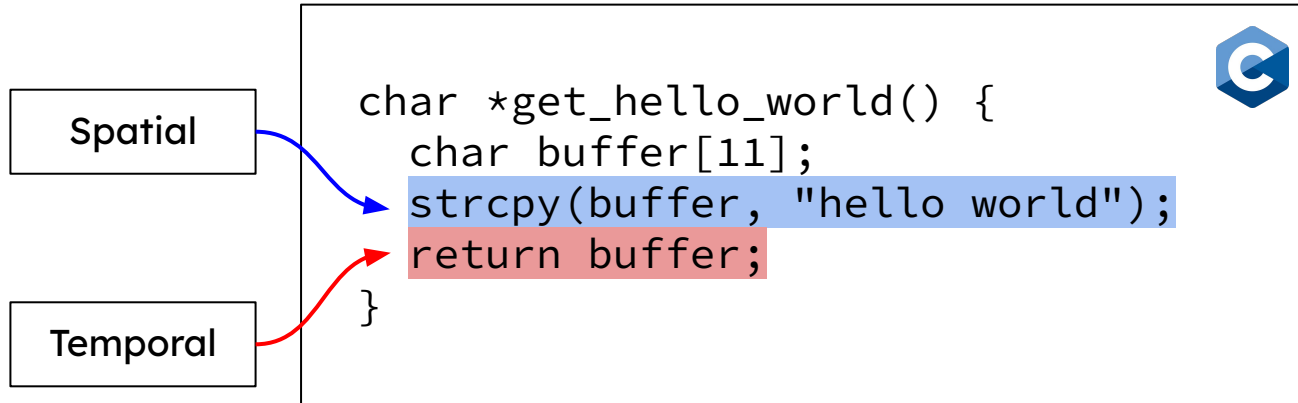
---

**Nearly 70%** of security bugs found by Google (2015 - 2020) and Microsoft (2006 - 2018) were caused by memory errors.

```
char *get_hello_world() {  
    char buffer[11];  
    strcpy(buffer, "hello world");  
    return buffer;  
}
```







# Garbage collection supports **dynamic memory safety**.

**Tracing garbage collection** treats memory as a reachability graph, and periodically eliminates nodes that are unreachable.


**Reference counting** frees memory when the count of references to an allocation in scope reaches zero.

---

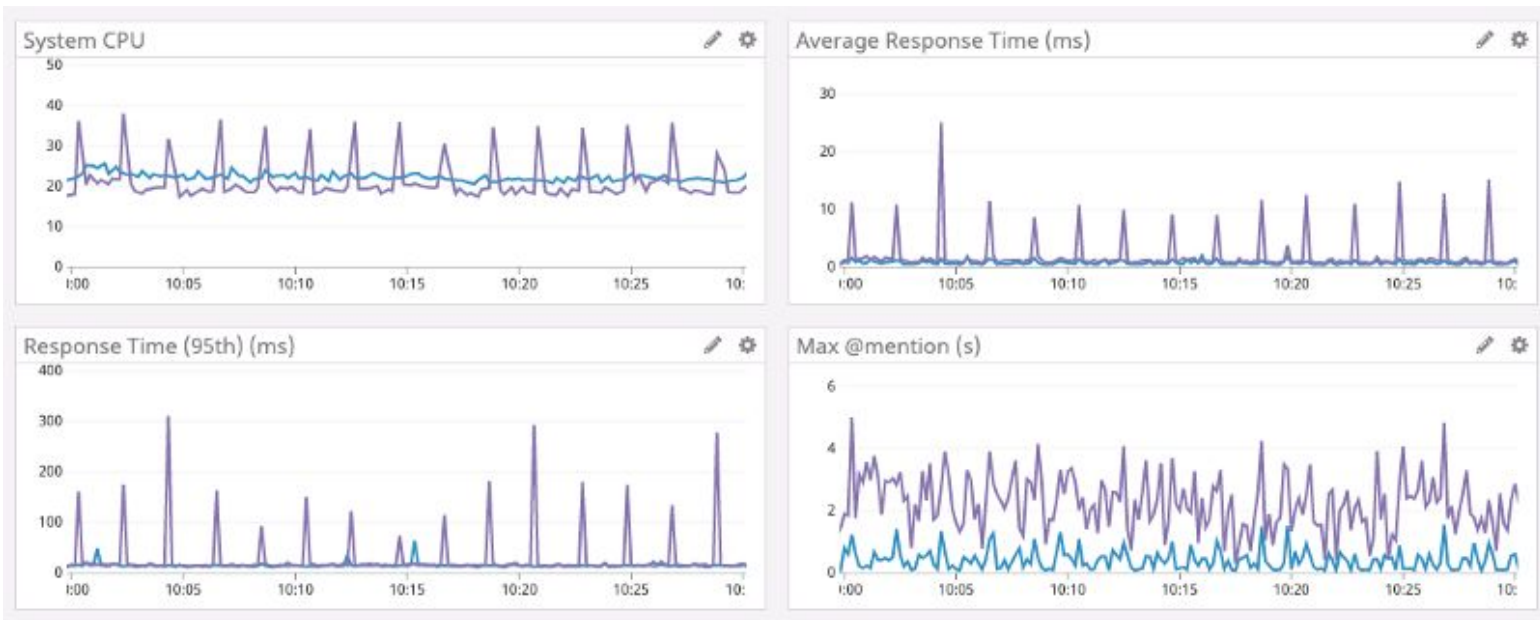
*No use-after free!*

- + No pointers
- + Bounds checks on array accesses

Check out *Crafting Interpreters*!



# Performance of Discord's “Read States” Service



Go Rust

(February 2020)

Rust is nearly as *performant* and *efficient* as C,  
but with **memory safety guarantees**.

# Rust

- Basic language features
- Ownership & the borrow checker
- Generic types

# Rust Resources

**The Rust Programming Language** — <https://doc.rust-lang.org/book/>

A tutorial on every aspect of Rust; a great starting point.

**The Rust Reference** — <https://doc.rust-lang.org/stable/reference/>

Comprehensive description of Rust's language features.

**Rust By Example** — <https://doc.rust-lang.org/rust-by-example/index.html>

Explanations of Rust concepts alongside runnable examples.

# Mutability

Variables are immutable by default.

Can be declared as mutable with the 'mut' keyword.

```
let x = 5;  
x = 7;
```



```
let mut x = 5;  
x = 7;
```

Why do we have  
immutability by default?

---

# Shadowing

```
let x:u16 = 0xffff;  
{  
    let x:u32 = 0xffffffff;  
    {  
        let x:u64 = 0xffffffffffffffff;  
    }  
}
```

```
let x:u16 = 0xffff;  
let x:u32 = 0xffffffff;  
let x:u64 = 0xffffffffffffffff;
```



# Numerical Types

Length	Signed	Unsigned
8-bit	<code>i8</code>	<code>u8</code>
16-bit	<code>i16</code>	<code>u16</code>
32-bit	<code>i32</code>	<code>u32</code>
64-bit	<code>i64</code>	<code>u64</code>
128-bit	<code>i128</code>	<code>u128</code>
arch	<code>isize</code>	<code>usize</code>

Number literals	Example
Decimal	<code>98_222</code>
Hex	<code>0xff</code>
Octal	<code>0o77</code>
Binary	<code>0b1111_0000</code>
Byte ( <code>u8</code> only)	<code>b'A'</code>

# Tuples & Arrays

```
let x:(u8, i32, char) = (1, -1000000, '@');  
  
let (byte, int, at) = x;  
  
let y = x.2;
```

```
let x:[i32;3] = [1, 1, 1];  
  
let x = [1;3];
```

# Functions & Closures

```
fn max(a: u32, b: u32) -> bool {  
  
    let min = || -> bool {  
        if a > b {  
            b  
        } else {  
            a  
        }  
    }  
  
    !min()  
  
}
```

# Structs & Traits

```
struct GradStudent {  
    coffee_allowance: u8,  
}  
  
trait Customer {  
    fn pay(& mut self, price: u8)  
}  
  
impl Customer for GradStudent {  
    fn pay(& mut self, price: u8) {  
        if self.coffee_allowance < price {  
            panic!("..  
        }else{  
            self.coffee_allowance -= price  
        }  
    }  
}
```

# The Borrow Checker

# The Borrow Checker

“...a system for statically building a proof that data in **memory is either uniquely owned** (and thus able to allow unguarded mutation) **or collectively shared, but not both.**”

Check out *Oxide: the  
Essence of Rust!*



# The Borrow Checker

1. All values have exactly one owner.

```
let x = String::from("Hello world!");
```

2. A reference to a value cannot outlive the owner.

```
let y = & mut x;
```

3. A value can have ***one mutable*** reference ***or many immutable*** references

You spend 3-6 months in a cave,  
*breathing, eating and sleeping* Rust.

Then find like-minded advocates

***who are prepared to sacrifice their first born***

to perpetuate the unfortunate sentiment that  
Rust is the future.

— Interviewee from [Fulton et al. 2021]




Learning Rust Ownership is

the

**best**

# The Borrow Checker rejects “valid” programs.




✗ Doubly-linked lists

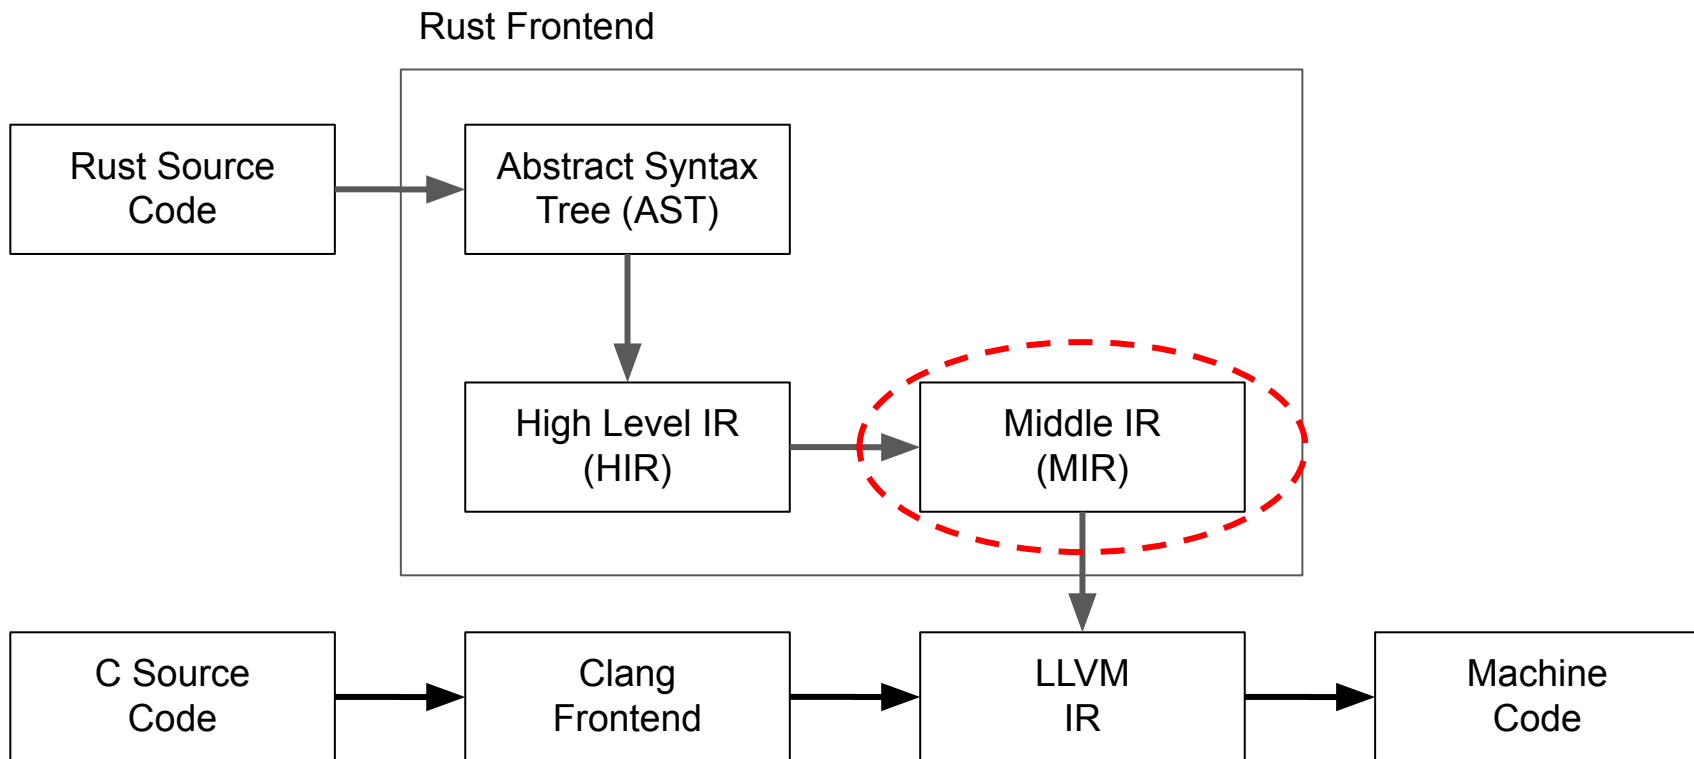
✗ Trees with parent and child pointers

✗ **Any** self-referential struct

# Self-Referential Structs


```
struct LLI {  
    context: Context,  
    module: Module,  
    engine: ExecutionEngine,  
}
```





### Copy Semantics

$x = y$




Y

X

### Move Semantics

$x = y$



Y

X

# Strings

```
let greeting = String::from("Hello world!");  
  
// create a response  
let response = greeting;  
  
// remove the '!'  
response.pop();  
  
// add a '?' to the end  
response.insert(response.len(), '?');  
  
println!("{}", greeting);  
  
println!("{}", response);
```




# Borrowing

```
let greeting: String = String::from("Hello world!");  
let response: &str = &greeting[..11];  
println!("{}", greeting);  
println!("{}", response);
```



# Slices

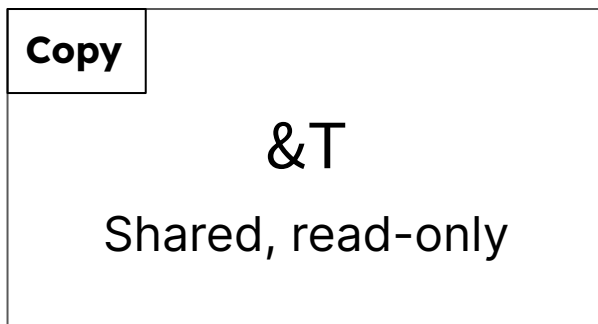




# Borrowing

```
let greeting: String = String::from("Hello world!");  
let response: &str = &greeting[..11];    // &[u8]  
println!("{}", greeting);  
println!("{}", response);
```





$\oplus$



# We can avoid moves with borrowing.

```
{  
    let mut greeting = String::from("Hello world!");  
  
    let response = &greeting[..11];  
  
    println!("{}", greeting);  
  
    greeting.clear();                // fn clear (& mut self)  
  
    println!("{}", response);  
}
```



# We can avoid moves with borrowing.

```
error[E0502]: cannot borrow `greeting` as mutable because it is also  
borrowed as immutable
```

```
--> src/lib.rs:8:2
```

```
4 |         let response = &greeting[..11];  
    |                               ----- immutable borrow occurs here  
...  
8 |         greeting.clear();  
    |         ^^^^^^^^^^^^^^^^^^^^^^^^^ mutable borrow occurs here  
9 |  
10 |         println!("{}", response);  
    |                               ----- immutable borrow later used here
```



Rust has awesome error messages!

Rust's borrow checker reasons using **lifetimes**.

The **scope of a value** is the duration for which it is **allocated**.

The **lifetime of a reference** is the *duration* for which it is **used**.

What are the scopes and lifetimes involved in this example?

```
{  
    let greeting = String::from("Hello world!");  
    let response = &greeting[..11];  
    println!("{}", greeting);  
    greeting.clear();    // implicit, & mut greeting  
    println!("{}", response);  
}
```



What are the scopes and lifetimes involved in this example?

```
{
    let greeting = String::from("Hello world!");
    let response = &greeting[..11];
    println!("{}", greeting);
    [
        greeting.clear();    // implicit, & mut greeting
        println!("{}", response);
    ]
}
```



Resource Allocation is Initialization (RAII)

```
#include <vector>

int main() {
    std::vector<int> v = {1, 2, 3, 4, 5};
    for (auto item : v) {
        v.push_back(item);
    }
}
```





```
fn main() {  
    let mut v = vec![1, 2, 3, 4, 5];  
    for item in v.iter() {  
        v.push(*item)  
    }  
}
```



error[E0502]: cannot borrow `v` as mutable because it is also borrowed as immutable  
--> src/main.rs:4:9

```
3 |     for item in v.iter() {  
    |                   -----  
    |                   |  
    |                   immutable borrow occurs here  
    |                   immutable borrow later used here  
4 |     v.push(*item)  
    |     ^^^^^^^^^^^^^^^^^ mutable borrow occurs here
```

# Generics

```
fn cashier<T:Customer>(customers: Vec<T>, items: Vec<u8> ) {  
    for (customer, item) in customers.iter().zip(items) {  
        customer.pay(item);  
    }  
}
```

# Enums

```
match something_error_prone() {  
    Ok(success_value) => { ... }  
    Err(err_value) => { ... }  
}
```

Result<T,E>

```
match optional_result() {  
    Some(value) => { ... }  
    None => { ... }  
}
```

Option<T>

# Shadowing

```
let value: Option<u64> = find_value();  
  
let value: u64 = match value {  
    Some(v) => v,  
    None => ...  
}
```

```
let value: u64 = find_value()  
    .unwrap_or(|| ...);
```

# Lifetimes

```
fn assign(mut a: & i32, b: & i32)
{
    a = b;
}
```



error: lifetime may not live long enough

[--> src/lib.rs:3:5](#)

```
1 | fn assign(mut a: & i32, b: & i32)
   |                               - let's call the lifetime of this reference `1`
   |                               |
   |                               let's call the lifetime of this reference `2`
2 | {
3 |     a = b;
   |     ^^^^^ assignment requires that `1` must outlive `2`
```

# Lifetimes

```
fn assign<'a, 'b>(mut a: & 'a i32, b: & 'b i32)
{
    a = b;
}
```



**'a :> 'b**

The region of code represented by **'b** must fully contain the region **'a**.

The region **'b** is a subtype of the region **'a**.

# Lifetimes are types.

```
class Animal { ... }  
class Crab extends Animal { ... }
```



```
Animal a = (Animal) new Crab();
```

1

```
Animal some_animal;  
...  
Crab a = (Crab) some_animal;
```

2



# Lifetimes

```
fn assign<'a, 'b>(mut a: & 'a i32, b: & 'b i32)
where 'b: 'a
{
    a = b;
}
```

**'a :> 'b**

The region of code represented by **'b** must fully contain the region **'a**.

The region **'b** is a subtype of the region **'a**.

# Lifetimes

```
fn assign<'a, 'b>(mut a: & 'a i32, b: & 'b i32)
where 'b: 'a
{
    a = b;
}
```

```
fn assign<'a>(mut a: & 'a i32, b: & 'a i32)
{
    a = b;
}
```

```
help: consider introducing a named lifetime parameter
1 | fn assign<'a>(mut a: &'a i32, b: &'a i32)
  |               +++++          ++          ++
```

## Exercise:

Complete the signature and implementation for the following versions of the 'swap' function without using any other dependencies. Some signatures are incomplete.

1. `fn swap(x: & i32, y: & i32)`
2. `fn swap(x: & & i32, y: & & i32)`
3. `fn swap<T>(x: & T, y: & T)`

Each should swap the values of the two parameters such that, after the function returns, we have:

`*y = old(*x) && *x = old(*y)`

# Answer #1

```
fn swap<T>(x: & mut i32, y: & mut i32)
{
    let temp = *x;
    *x = *y;
    *y = temp;
}
```

## Answer #2

```
fn swap<T, 'a>(x: & mut 'a & T, y: & mut 'a & T)
{
    let temp = *x;
    *x = *y;
    *y = temp;
}
```

## Answer #3

```
fn swap<T>(x: & mut T, y: & mut T)
{
    unsafe {
        let a = ptr::read(x);
        let b = ptr::read(y);
        ptr::write(x, b);
        ptr::write(y, a);
    }
}
```

`std::mem::swap`